

MongoDB: Distributed NoSQL database

Martin Carrasco
Computer Science
UTEC
Lima, Peru
martin.carrasco@utec.edu.pe

Sebastian Carcamo
Computer Science
UTEC
Lima, Peru
sebastian.carcamo@utec.edu.pe

Jorge Ignacio Fiestas
Computer Science
UTEC
Lima, Peru
jorge.fiestas@utec.edu.pe

I. INTRODUCTION

A. Description & Motivation

MongoDB was originally developed as a component for a *platform as a service* product. These were web based application that looked to simplify the development and deployment process of web based app for developers. For this reason, *MongoDB* was conceived with two main purposes: 1) Being flexible and simple to use, as developers might need to use it for a wide range of applications; and 2) Being scalable in order to allow to work with hughmongous amounts of data (that is were it gets its name).

MongoDB is a document based non-relational database model that works with three main components: documents, fields and collections. *Documents* would be the *MongoDB* equivalent of a register in a traditional SQL database. Each documents has *fields*, that act like the attributes or columns of each register. This documents stored in BSON, which is a JSON file that is stored as a binary to optimize memory. When we group some documents we get a *collection*, that are the closest thing *MongoDB* has to tables and schemas.

One of the key differences between *MongoDB* and relational database models is that documents in the same collection do not need to share structure. That is to say that we can have documents with different fields and that not all objects with the same field have to store the same type of data. This also allows for the stored data to have more semantic meaning, and therefore, be more *human readable*.

B. Comparison with other Non-Relational DB

MongoDB is currently established as the most popular non-relational database model in the market. This is probably due to three main differentiating factors:

- 1) As *MongoDB* was one of the few DB developed in C++ it doesn't have the overhead that some of its competitors have, allowing it to be fast and lean.
- 2) *MongoDB* data storage is based on the filesystem of the native OS of the system in which it is working, this make it portable and lightweight, as it doesn't depend on external filesystems or further abstraction layers.
- 3) Even though it doesn't have its own query language, there exist multiple *MongoDB* libraries for almost all programming languages, specially web-oriented ones.

DB/ Properties	MongoDB	Cassandra	Accumulo	Couch DB	Hbase	Redis	Riak
Language	C++	Java	Java	Erlang	Java	C.C++	Erlang
Data Model	BSON	Big Table	Big Table	JSON	Big Table and Dynamo	Data Structure	Data structure
Fault Tolerance	Replication	Partitioning and replication	Replication	Replication	Partitioning and replication	Replication	Replication
Data Storage	Memory, file framework	Dynamo for storing data	HDFS	Memory, File framework	HDFS	File system	Bit cast, Memory
Community	AGPL	Facebook	Apache	Apache	Apache	BSD	Apache
MapReduce	YES	YES	Yes	YES	YES	NO	YES
Query Language	-	API calls	Java API, Thrift API	-	XML, Thrift API	API calls	JavaScript
Replication Modes	Master-Slave Replication	Master-Slave replication	Multi-master replication	Multi-Master Replication	Master-Slave Replication	Master-Slave Replication	Multi Master Replication
Protocol	TCP/IP	Thrift	Thrift	HTTP/ REST	Thrift, API, tradition	Binary, Similar to telnet	REST

Fig. 1. A comparison between different non-SQL DB Models.

C. Use Cases & Applications

MongoDB main use case is in the development of web based applications that need to evolve quickly and scale elegantly. This makes sense as it's JSON based file storage is highly compatible with most front-end programming languages. As well, thanks to it's simple file structure, *MongoDB* is good dealing with multidimensional data such as text, geospatial data and time series. As there is no restrictions to which documents get grouped in a collection, *MongoDB* is also perfect for the representation of data with natural clusters or high variability.

Thanks to its flexibility and reliability, *MongoDB* is used by tech leads businesses such as Adobe, Google, Nokia, SAP and many more.

II. CHARACTERISTICS DESCRIPTION

In this section, we will talk about the functionalities that *MongoDB* has to offer. In this particular case, we will describe specifically about the *MongoDB* shell method instead of using APIs since this will give us more of an inside look of what's really happening in detail.

A. Data model

First of all, its worth mentioning that *MongoDB* gives the developers a lot of control over something called *durability* in their write and read calls. Said durability is a property that guarantees that the written data will be saved and be permanent. This leads to a decision to be taken by the developer as a strong durability parameter will be safer to use since its guaranteed that the data will be saved on disk whilst

a weak durability parameter will lead to a higher throughput but provide less data safety.

1) *Insertion*: The insertion method is used to insert one or more documents to a collection. Said documents need to be tagged with an id, whether its given by the user or not. If the user doesn't provide a unique *_id* field the daemon will assign it one.

```
db.products.insert( { item: "card",  
  ↪ qty: 15 } )  
  
{ "_id" :  
  ↪ ObjectId("5063114bd386d8fadbd6b004"),  
  ↪ "item" : "card", "qty" : 15 }
```

Listing 1: Input and output of an insertion without specifying *_id*

```
db.products.insert( { _id: 10, item:  
  ↪ "box", qty: 20 } )  
  
{ "_id" : 10, "item" : "box", "qty" :  
  ↪ 20 }
```

Listing 2: Input and output of an insertion specifying *_id*

Multiple documents can be inserted if you pass an array of documents as the argument to the insert() method.

```
db.products.insert(  
  [  
    { _id: 11, item: "pencil", qty:  
      ↪ 50, type: "no.2" },  
    { item: "pen", qty: 20 },  
    { item: "eraser", qty: 25 }  
  ]  
)  
  
{ "_id" : 11, "item" : "pencil", "qty"  
  ↪ : 50, "type" : "no.2" }  
{ "_id" :  
  ↪ ObjectId("51e0373c6f35bd826f47e9a0"),  
  ↪ "item" : "pen", "qty" : 20 }  
{ "_id" :  
  ↪ ObjectId("51e0373c6f35bd826f47e9a1"),  
  ↪ "item" : "eraser", "qty" : 25 }
```

Listing 3: Input and output of an insertion of multiple documents

Regarding the starting introduction to durability, we have the *writeConcern* parameter. We will talk more in detail about this parameter later.

Ordered is another interesting parameter that MongoDB gives us choice to modify. When inserting in bulk, the default value for ordered is true. Doing an ordered insert implies that the insertions will be done serially and if an error occurs during one of those insertions, MongoDB will return without processing any of the remaining write operations in the list. On the other hand, an unordered insertion writes in parallel and will not stop writing if one of the insertions fails.

Lastly, its worth mentioning that the insertOne() and insertMany() methods are preferred since what they return gives more details about the insertions made.

2) *Deletion*: The deletion is done with the remove() method. This method takes the deletion criteria as a *query* parameter. It also has a parameter *justOne* that when set to true will only delete the first element that it finds that matches the query. Similarly to the insert method, the remove method also has a *writeConcern* parameter. The last parameter is *collation*, this parameter allows users to specify language-specific rules for string comparison.

Again, the alternative deleteOne() and deleteMany() methods are preferred because of their return values.

3) *Updates*: Updates can be performed in MongoDB collections. To do so, we make use of the updateOne() and updateMany() methods. Both methods work similarly. The first argument that they take is the condition that the document must meet to be modified. The second argument is the list of operations to be done to those documents.

```
db.collection.update(  
  <query>,  
  <update>,  
  {  
    upsert: <boolean>,  
    multi: <boolean>,  
    writeConcern: <document>,  
    collation: <document>,  
    arrayFilters: [ <filterdocument1>,  
      ↪ ... ],  
    hint: <document|string>  
  }  
)
```

Listing 4: Update() method

A list with the update operators that can be applied to the documents can be found in the [MongoDB documentation](#)

4) *BulkWrite*: All the previous methods don't have to be written one by one, they can all be performed inside a bigger method called bulkWrite(). This method takes all the operations and uses the specified *writeConcern* and *ordered* parameters and uses it for all of them.

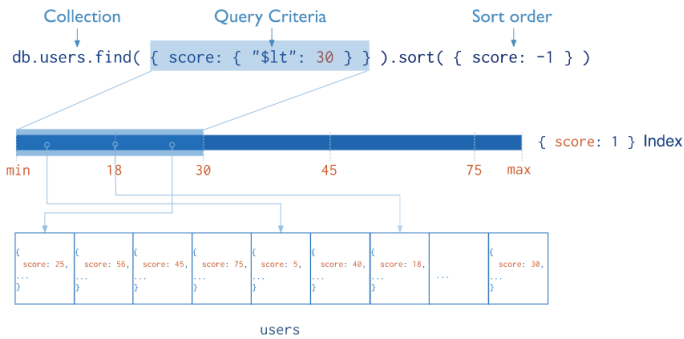


Fig. 2. Caption

```
db.collection.bulkWrite(
  [ <operation 1>, <operation 2>, ...
    ↪ ],
  {
    writeConcern : <document>,
    ordered : <boolean>
  }
)
```

Listing 5: bulkWrite() method

5) *Search*: To select documents in a collection or view we use the find() method. The first parameter that it receives is the query, where you can use any of these [MongoDB query operators](#). The second parameter specifies the fields of the matching documents to be shown. Both of these parameters are optional and in case the user runs the method without parameters, the result will be all the documents in the collection.

6) *Indexing*: Indexes support the efficient execution of queries in MongoDB. Without them, a scan in every single document in the collection should be done to find the documents that match a query, which is very inefficient. Indexing limits the amount of documents MongoDB needs to inspect. MongoDB uses the B-tree data structure to allow the users to create indexes that store a small portion of the collection's data set in an easy to traverse form.

MongoDB creates a unique index on the _id field when creating a collection and also allows the users to create indexes based on their selected fields.

```
db.collection.createIndex( <key and
  ↪ index type specification>,
  ↪ <options> )
```

Listing 6: Command to create an index on a collection

III. DISTRIBUTED STORAGE ARCHITECTURE

A. Sharding

Given a data set too large for a single server to handle or high throughput applications we face a number of problems:




Fig. 3. Range-based sharding

high query rates exhaust the CPU capabilities of a server. Working set sizes larger than the server's RAM capacity stress the I/O capacity of disk drives.

To solve these problems we need to address the two types of scaling to handle system growth: vertical and horizontal scaling.

Vertical scaling consists in increasing the capabilities of a single server so its basically giving it more RAM or increasing the storage space. This type of scaling is limited by the current available technology.

Horizontal scaling consists in dividing the workload between multiple servers. The speed or capacity of a single server may not be high but while it works together with more servers, the load on a single server is diminished. Nonetheless, the trade-off for gaining efficiency is adding complexity to the infrastructure and maintenance to the deployment.

Adding this horizontal scaling in MongoDB is referred as *sharding*. To do this, we need to specify a shard key. A shard key is a field that must exist in every document, that will be used to distribute the documents across shards in chunks. Said chunks have inclusive lower and exclusive upper range based on the shard key.

There are two types of sharding that can be used within MongoDB: Range-based sharding and Hash-based sharding. The range-based sharding divides the data sets into different ranges based on the value of the shard key. Documents having close shard key values are more likely to be in the same shard. This type of sharding enables us to do range queries based on the shard key values since its very probable that the adjacent document is located in the same shard. However, this partitioning can be uneven and will fail to scale if a small set of shards receive most of the requests.

On the other hand, the hash-based sharding is done by calculating the hash of the document's shard key value and then distributing the documents according to the result. This results in an evenly distributed data among the shards but doesn't guarantee that the documents with close shard key values reside in the same shard. This renders ineffective the range queries.

A MongoDB sharded cluster consists of the following components:

shards: contain subsets of the data set. mongos: Act as query router and provide an interface between client applications and the sharded cluster config servers: store metadata and configuration settings.




Fig. 4. Hash-based sharding




Fig. 5. Components

B. Considerations before sharding

Sharding a data set requires infrastructure setup, which increases the deployment complexity. Consider sharding clusters when these apply:

- The data set is larger than a single MongoDB storage capacity.
- The size of the active working set exceeds the capacity of the maximum available RAM.
- A single MongoDB instance is unable to manage write operations.

Also, a pertinent choice of shard key is needed for ensuring efficiency and maximize performance as this key cannot be changed later.

C. Replication

A replica set is a group of mongod processes that maintain the same data set. Having replica sets provides redundancy and high availability of the data when needed.

The sequence for writing data considering the replica sets is the following:

- 1) The write operations go to the primary node
- 2) The changes are recorded into the operation log
- 3) The write operations are replicated to the secondary nodes
- 4) Secondaries copy the primary oplog

The write operations on the secondary nodes and copying the primary oplog are done asynchronously. Having the replica set setup like this allows us to continue to function despite one or more nodes to be down or have failed.

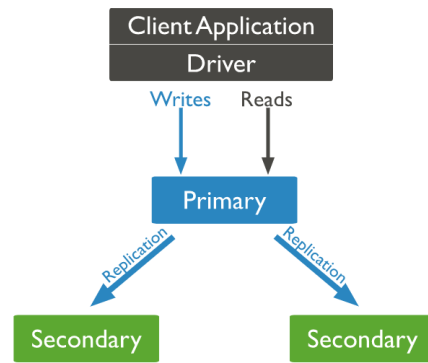


Fig. 6. Connections between nodes

When the primary node does not communicate with the secondaries after a certain amount of time, an election for a new primary node is done within all the nodes so a new primary is assigned and the functions can continue. The replica set cannot process write operations until the election completes successfully. The replica set can continue to serve read queries if such queries are configured to run on secondaries while the primary is offline.

The configuration options for the replica sets are:

- 50 members per replica set (7 voting members)
- Arbiter node:
 - Does not hold a copy of the data
 - Only used to vote in the elections to form a majority
- Priority 0 node:
 - Never becomes primary
 - Visible to the application
- Hidden node:
 - Never becomes primary
 - Not visible to the application
 - Used for reporting or for backups
- Delayed node:
 - Has a delay for a certain amount of seconds
 - Never becomes primary
 - (Optional) Hidden
 - Votes in elections
 - Mainly used for backups
- Write concern:
 - w - number of mongod instances (replica members) that acknowledged the write operation
 - j - number of mongod instances that acknowledged that the write operation has been written to their respective journal
 - wtimeout - time limit to prevent write operations
- Read preference:
 - primary (default) - reads only from the primary node
 - primaryPreferred - reads from the primary as a default but will read from the secondaries in case of failure.

- secondary - reads only from a secondary node
- secondaryPreferred - reads from the secondaries as a default but will read from the primary in case of failure.
- nearest - will read from the node that has the least network latency

IV. IMPLEMENTATION

A. Dataset information

We used the dataset provided as one of the examples from MongoDB. This is the *Amsterdam AirBnB* dataset. It contains around 19K entries for AirBnB offers. Also taken from the MongoDB page here is an extract with a sample document general properties. The values of the fields are omitted to preserve the order in the document. We inserted this entries and indexed on the field *id* and *host_since*.

```
{
  "_id": ,
  "listing_url": ,
  "name": ,
  "summary": ,
  "interaction": ,
  "house_rules": ,
  "property_type": ,
  "room_type": ,
  "bed_type": ,
  "minimum_nights": ,
  "maximum_nights": ,
  "cancellation_policy": ,
  "last_scraped": ,
  "calendar_last_scraped": ,
  "first_review": ,
  "last_review": ,
  "accommodates": ,
  "bedrooms": ,
  "beds": ,
  "number_of_reviews": ,
  "bathrooms": ,
  "amenities": ,
  "price": ,
  "security_deposit": ,
  "cleaning_fee": ,
  "extra_people": ,
  "guests_included": ,
  "images": ,
  "host": ,
  "address": ,
  "availability": ,
  "review_scores": ,
  "reviews":
}
```

Listing 1. Document sample

B. Setup and configuration of MongoDB

1) *Basic setup*: The configuration we used is the following. We have 3 configuration servers, 3 shards with a replica each and a router. The *docker_compose.yml* used to deploy the services are at the Appendix as well as the scripts for the configuration of each shard, the router and the configuration server. We will also reference the scripts used in our compose. Essentially we follow the tutorial referenced in the bibliography and sharded our collection according to some fields. To test the queries and load the dataset into MongoDB we decided to use Mongo Compass, a GUI that allows us to perform queries, upload data and, most importantly, analyze our queries with graphs to be able to better understand what is going on behind the curtains. Now we proceed to explain the steps we took next, which will be done in Mongo Compass, since all the configuration is already in place (except for the creation of the index and sharding) for us to start using MongoDB. In Figure 7 we can see the configuration using Mongo Compass about our sharding setup. We can also see in Figure 8 how our data is represented in Mongo Compass. We will later use some functions like filtering to run queries on our data and do the explain analyze of what happened with these queries.

2) *Index and collection sharding*: To test the performance of our database we will run the same 3 queries in 2 setups. First we use a composite index of 3 fields *bedroom_type*, *room_type* and *property_type*. This one has no sharding. For the second setup we use a hash index on each of the fore mentioned fields and performed a sharding of the collection on *room_type*.

C. Results and performance

To test our MongoDB setup we loaded the data and performed indexing as said before previously and the sharding. We will analyze the 2 setups. We will explain the queries and show the outputs of the explain analyze as follows.

1) Setup 1: No sharding:

- Here we filter all AirBnB listings that contain 1 bedroom and 2 beds and sort by the *host_since* date field and limit results at 1000. The total execution time is 114ms and the execution time performed in the shard01 is 54ms. We can see the results in Figure 9
- Here we perform an exact match to search for all the hosts named *Daniel* and we don't cap the results. We see here that the query takes 32ms and the time in the shard is 30ms. A reasonable time with 44 returned documents. We can see the results in 10

2) Setup 2: Sharding on *room_type*:

- Here we perform the same query as the first part, however we see that the total time is 204ms. This is due to gathering from all shards and then performing a merge. Since the query is not on the sharded attribute it is expected to be slower. 11
- We perform an exact match on a host name and the sharding actually increases the time from no sharding to 40ms but just a little. 12
- Finally we perform a query on one of the sharded attribute and sort by host since. This takes a time of of 100ms

which is way lower than the time it takes to run without sharding. This proves we need to be carefull about the indexing and sharding of these related to the type of work we want to do. 13

V. CONCLUSION

After learning about how MongoDB performs in a distributed enviroment and how NoSQL data is handled by DBMS we can reach some final conclusions as to what is the value of MongoDB is the database ecosystem.

REFERENCES

- [1] Veronika Abramova and Jorge Bernardino. Nosql databases: Mongodb vs cassandra. In *Proceedings of the international C* conference on computer science and software engineering*, pages 14–22, 2013.
- [2] Alexandru Boicea, Florin Radulescu, and Laura Ioana Agapin. Mongodb vs oracle–database comparison. In *2012 third international conference on emerging intelligent data and web technologies*, pages 330–335. IEEE, 2012.
- [3] Kristina Chodorow. *MongoDB: the definitive guide: powerful and scalable data storage*. ” O’Reilly Media, Inc.”, 2013.
- [4] Peter Membrey, Eelco Plugge, and DUPTim Hawkins. *The definitive guide to MongoDB: the noSQL database for cloud and desktop computing*. Apress, 2011.
- [5] Thomas Stevens. mongo-shard-docker-compose. <https://github.com/chefspate/mongo-shard-docker-compose>, 2018.

APPENDIX

```
{
version: '2'
services:
  config01:
    image: mongo
    command: mongod --port 27017 --
      ↪ configsvr --replSet configserver
    volumes:
      - ./scripts:/scripts
  config02:
    image: mongo
    command: mongod --port 27017 --
      ↪ configsvr --replSet configserver
    volumes:
      - ./scripts:/scripts
  config03:
    image: mongo
    command: mongod --port 27017 --
      ↪ configsvr --replSet configserver
    volumes:
      - ./scripts:/scripts
  shard01a:
    image: mongo
    command: mongod --port 27018 --shardsvr
      ↪ --replSet shard01
    volumes:
      - ./scripts:/scripts
  shard01b:
    image: mongo
    command: mongod --port 27018 --shardsvr
      ↪ --replSet shard01
```

```
volumes:
  - ./scripts:/scripts
shard02a:
  image: mongo
  command: mongod --port 27019 --shardsvr
    ↪ --replSet shard02
  volumes:
    - ./scripts:/scripts
shard02b:
  image: mongo
  command: mongod --port 27019 --shardsvr
    ↪ --replSet shard02
  volumes:
    - ./scripts:/scripts
shard03a:
  image: mongo
  command: mongod --port 27020 --shardsvr
    ↪ --replSet shard03
  volumes:
    - ./scripts:/scripts
shard03b:
  image: mongo
  command: mongod --port 27020 --shardsvr
    ↪ --replSet shard03
  volumes:
    - ./scripts:/scripts
router:
  image: mongo
  command: mongos --port 27017 --configdb
    ↪ configserver/config01:27017,
    ↪ config02:27017,config03:27017 --
    ↪ bind_ip_all
  ports:
    - "27017:27017"
  volumes:
    - ./scripts:/scripts
  depends_on:
    - config01
    - config02
    - config03
    - shard01a
    - shard01b
    - shard02a
    - shard02b
    - shard03a
    - shard03b
```

Listing 2. docker_compose.yml

```
rs.initiate(
{
  _id: "configserver",
  configsvr: true,
  version: 1,
  members: [
    { _id: 0, host : "config01:27017"
```

```

        ↪ },
    { _id: 1, host : "config02:27017"
        ↪ },
    { _id: 2, host : "config03:27017" }
  ]
}
)

```

Listing 3. init-configserver.js

```

    { _id: 1, host : "shard03b:27020"
        ↪ },
  ]
}
)

```

Listing 7. init-shard03.js

```

sh.addShard("shard01/shard01a:27018")
sh.addShard("shard01/shard01b:27018")

sh.addShard("shard02/shard02a:27019")
sh.addShard("shard02/shard02b:27019")

sh.addShard("shard03/shard03a:27020")
sh.addShard("shard03/shard03b:27020")

```

Listing 4. init-router.js

```

rs.initiate(
{
  _id: "shard01",
  version: 1,
  members: [
    { _id: 0, host : "shard01a:27018"
        ↪ },
    { _id: 1, host : "shard01b:27018"
        ↪ },
  ]
}
)

```

Listing 5. init-shard01.js

```

rs.initiate(
{
  _id: "shard02",
  version: 1,
  members: [
    { _id: 0, host : "shard02a:27019"
        ↪ },
    { _id: 1, host : "shard02b:27019"
        ↪ },
  ]
}
)

```

Listing 6. init-shard02.js

```

rs.initiate(
{
  _id: "shard03",
  version: 1,
  members: [
    { _id: 0, host : "shard03a:27020"
        ↪ },
  ]
}
)

```

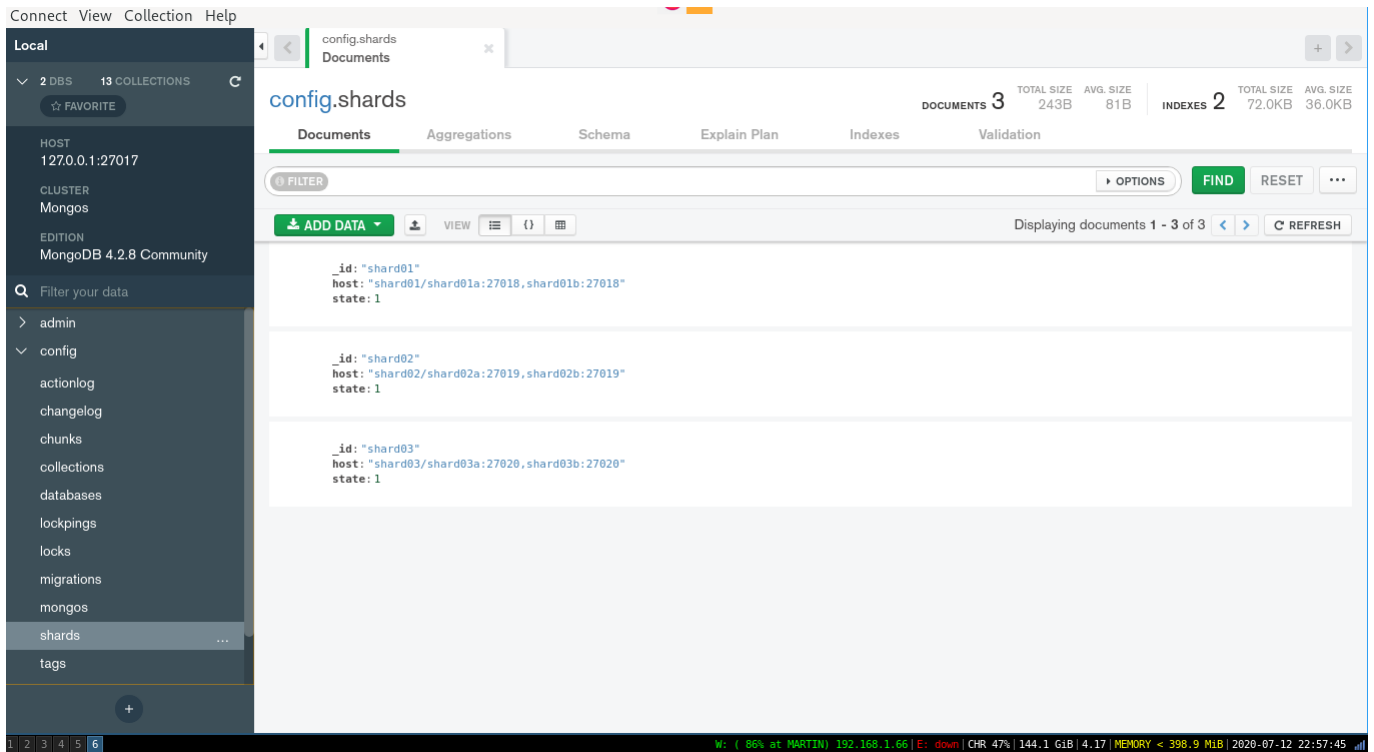


Fig. 7. Sharding Setup

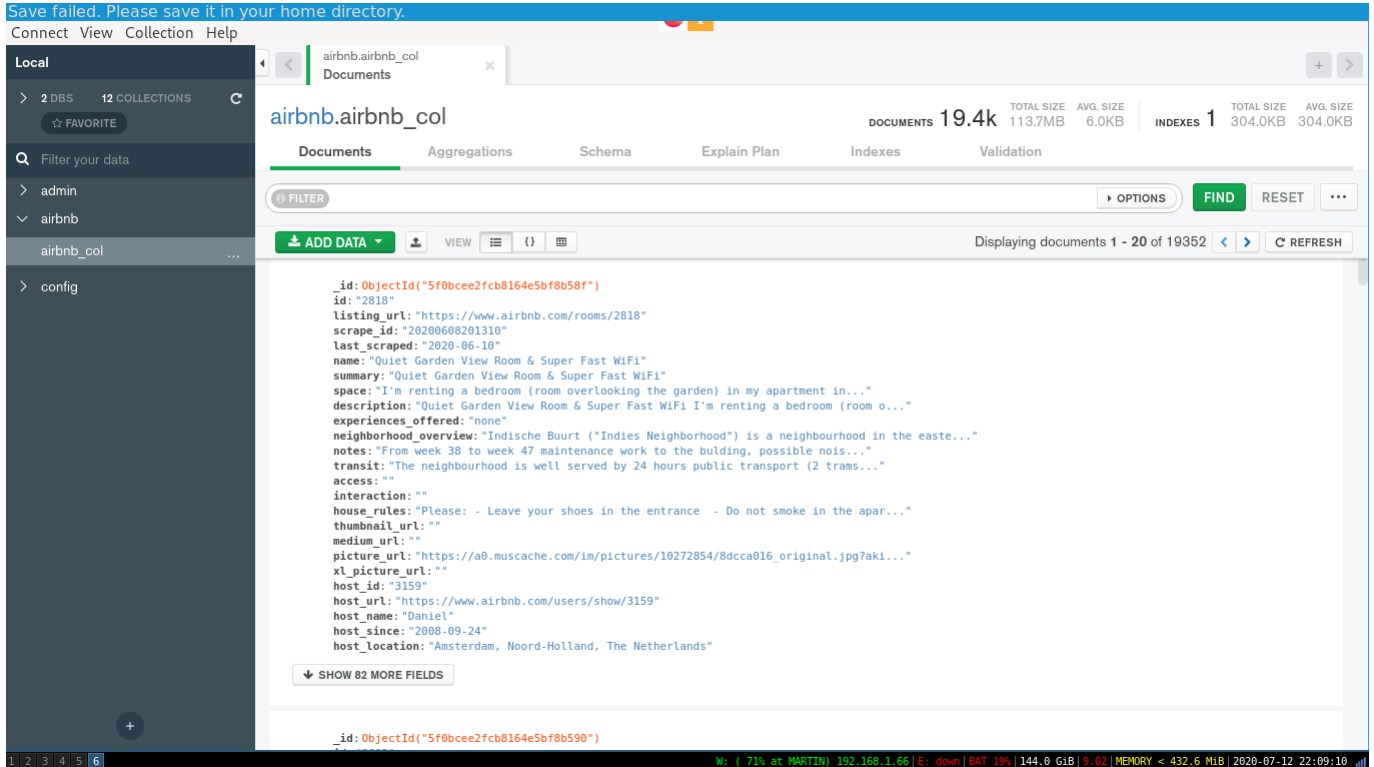


Fig. 8. Mongo Compass General View

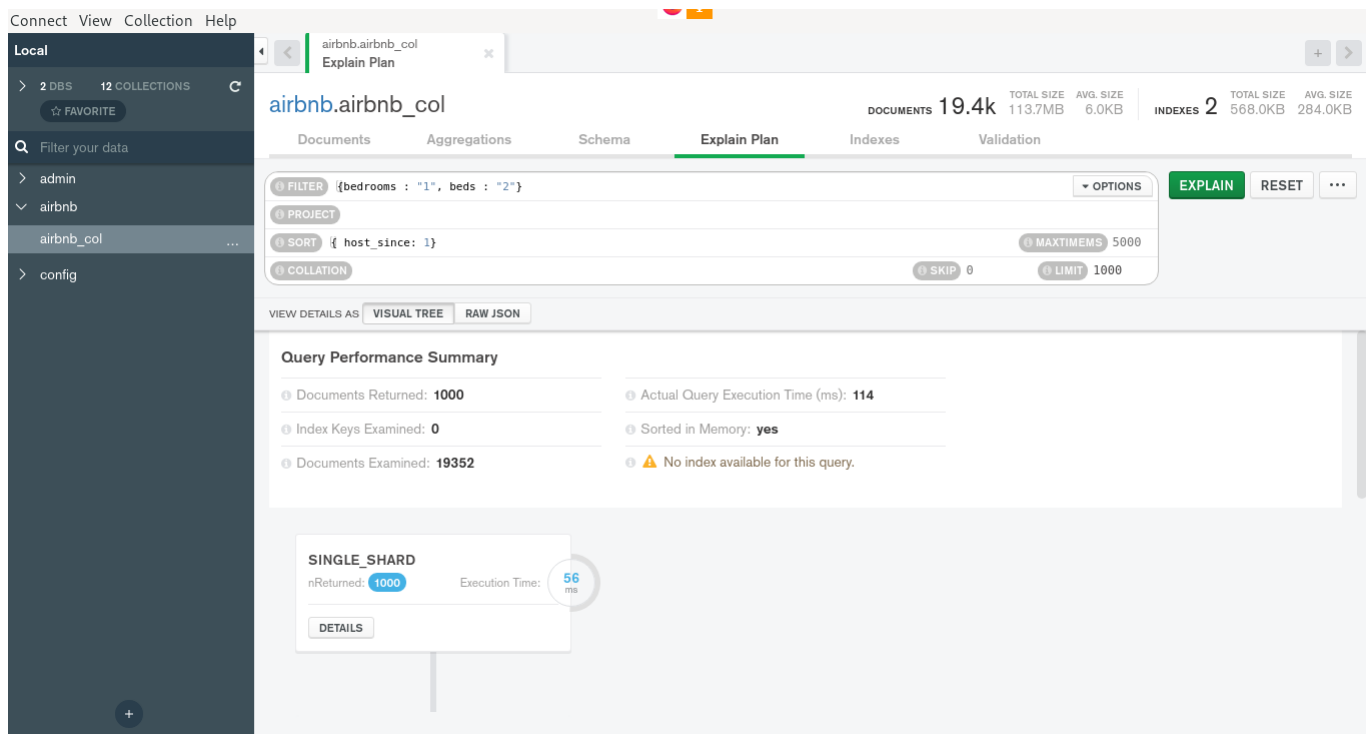


Fig. 9. Query 1

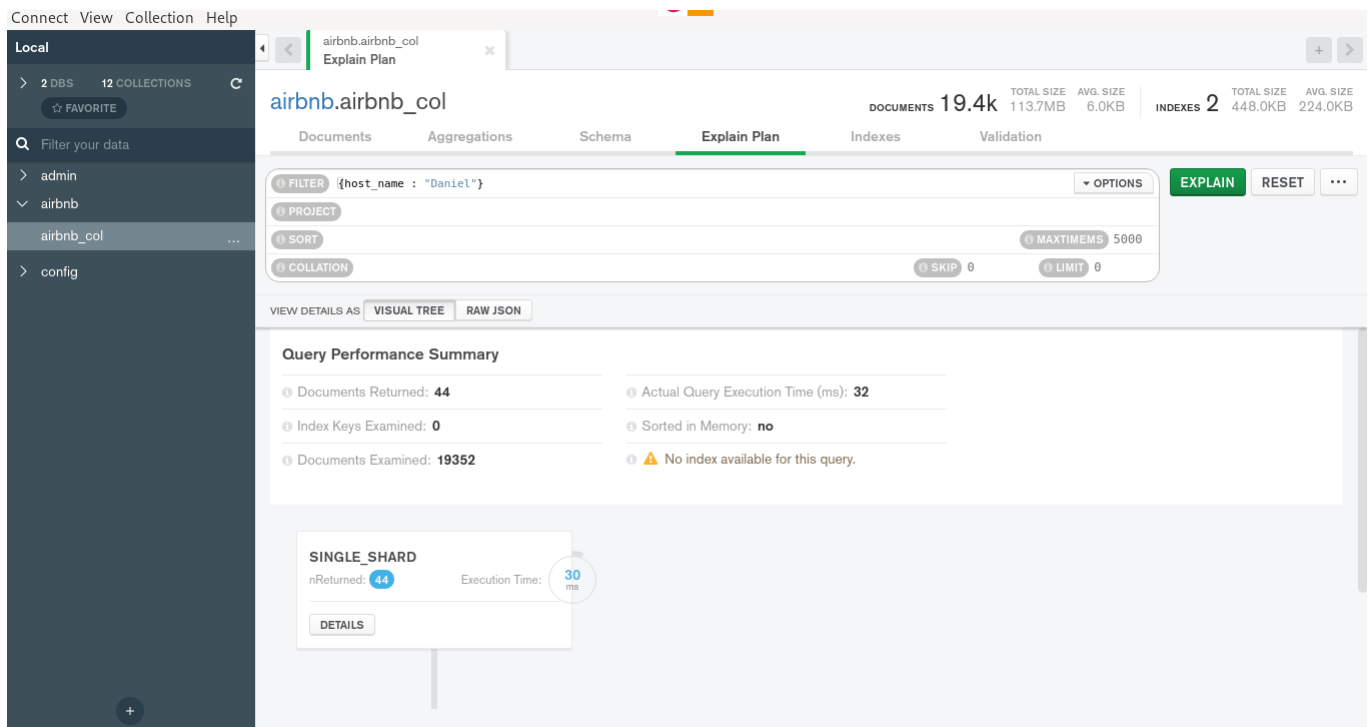


Fig. 10. Query 2

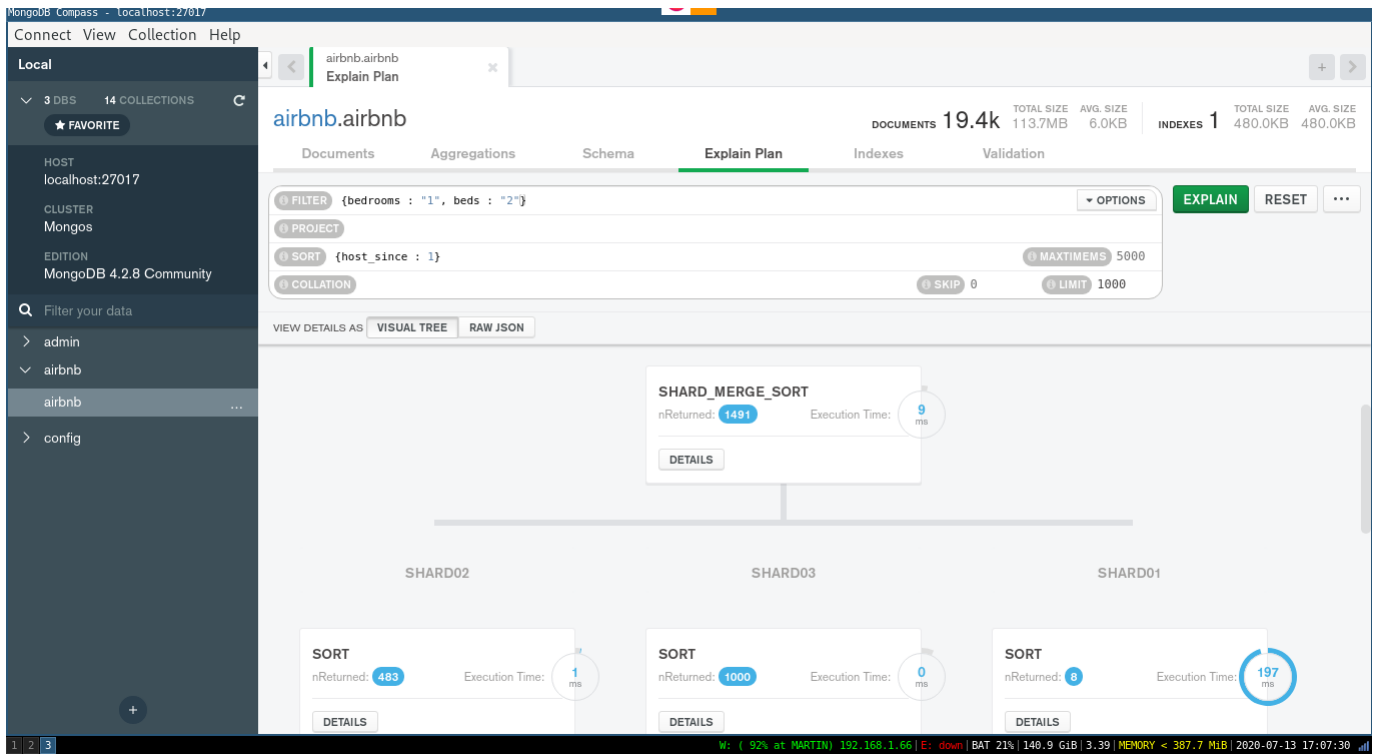


Fig. 11. Query Sharded 1

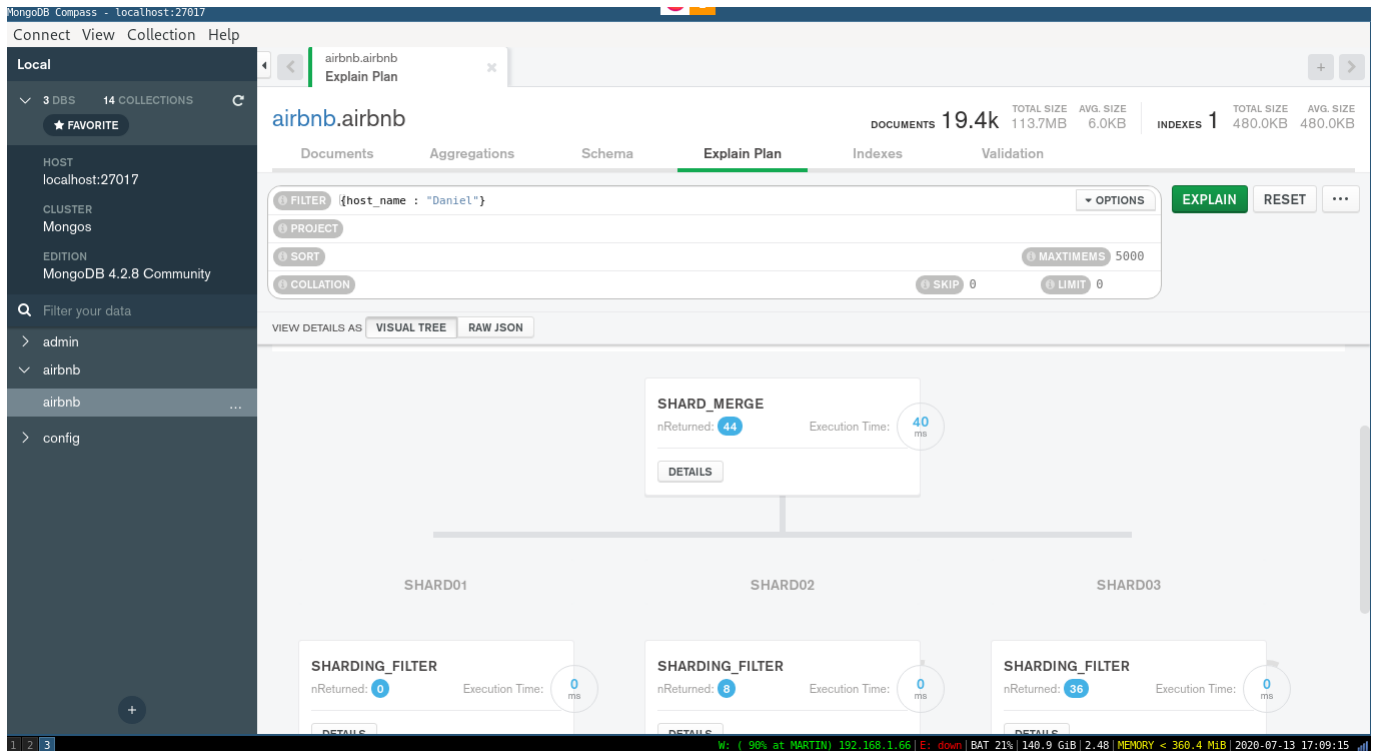


Fig. 12. Query Sharded 2

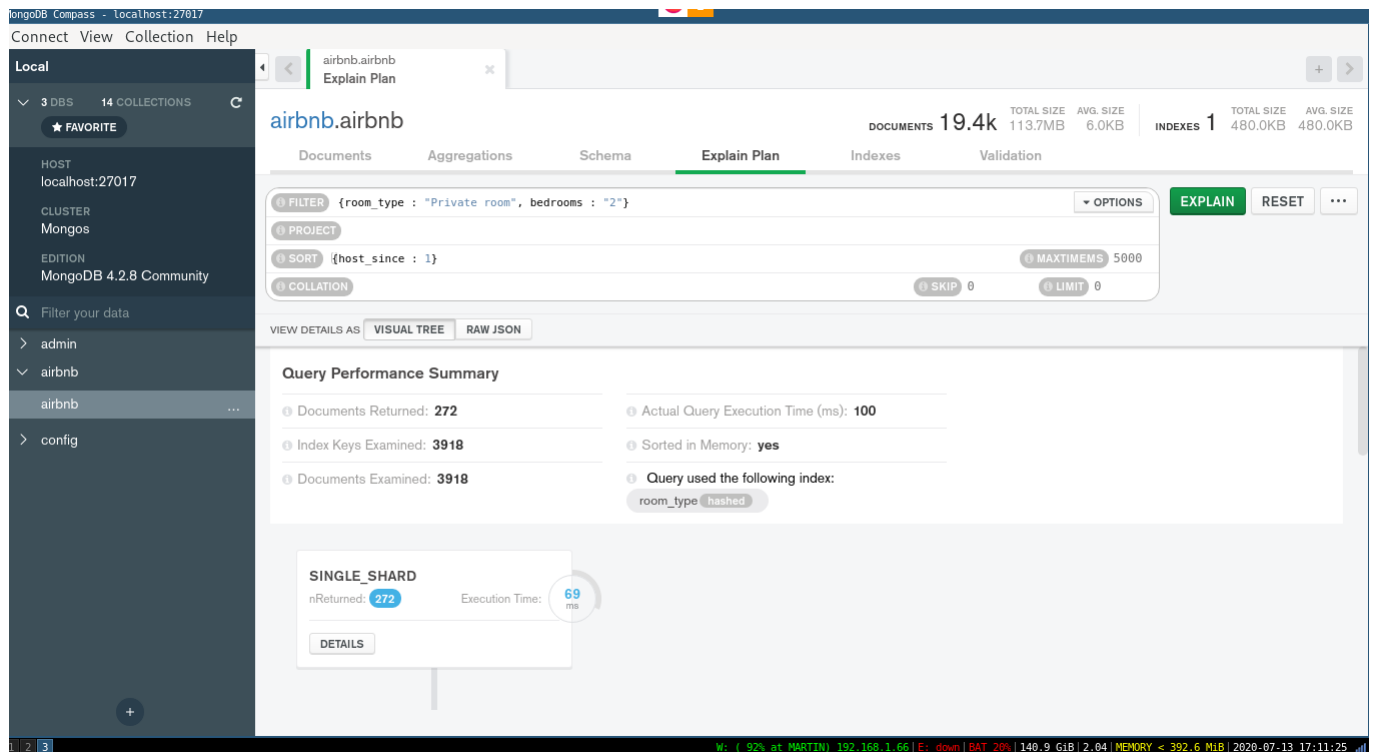


Fig. 13. Query Sharded 3